CSCI E-82a
Probabilistic Programming and AI
Lecture 1
Directed Graphical Models

Steve Elston

Why Directed Graphical Models?

How can we make inferences in an uncertain environment?

- Example; **sensor fusion** for a self-driving car
- Example; Locate robot while mapping environment
- Example; make a decision given **uncertain information**
- Need compute the **posterior of multivariate probability distribution**:

$$p(x_1, x_2, \ldots x_n)$$

- Or, compute **maximum a posteriori** (MAP)
- Process of **inference**

Why Directed Graphical Models?

Are directed graphical models interpretable?

- Allow experts to include **prior information**!

- Provide a language to describe the relationship between variables

- Can easily query **conditional probability tables** (CPTs)

- But, can non-experts really understand these models?
  - Making the model results transparent to non-experts is a research area

Why Directed Graphical Models?

We want the **posterior distribution** of one or more variables

- How to compute this **query** efficiently?

- Direct tabular representation **O(n²) complexity**

- Can do (much) better **factoring distribution by conditional independencies**

  - Discuss complexity of inference later

- **Directed graphical model is a representation of a distribution factored by conditional independencies**

Why Directed Graphical Models?

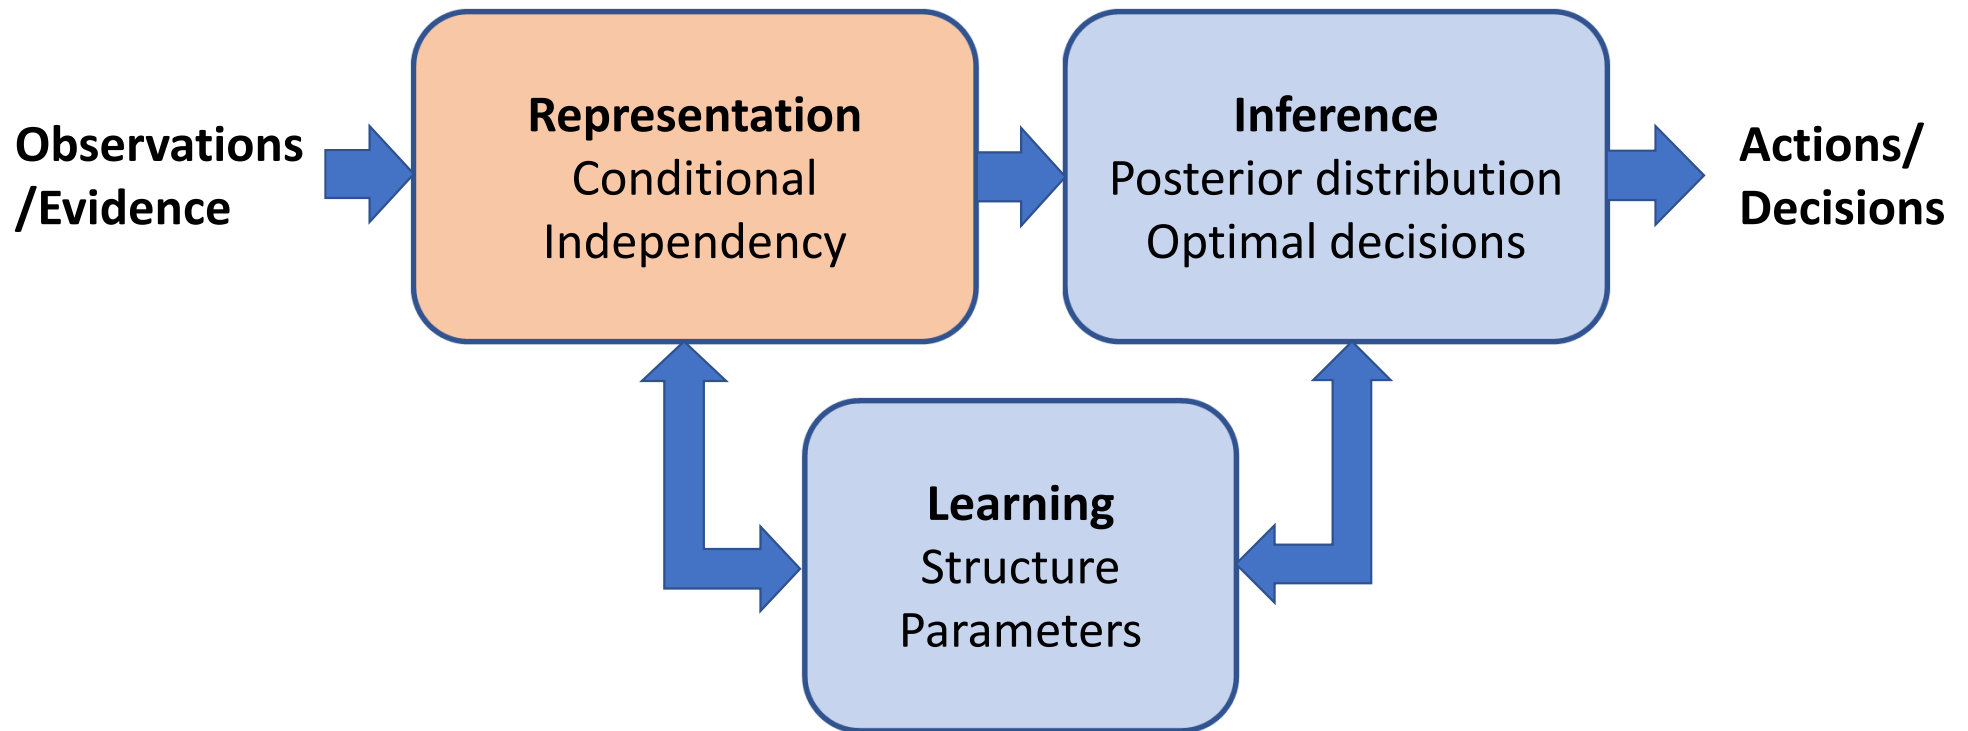Directed graphical models are generative

- Define probability distribution with DAG

- DAG is expressive in terms of independency structure

- Generate realizations from that distribution

Why Directed Graphical Models?

How much data do we need for learning?

- Learning is both hard and easy
- Model parameters (CPTs) improve with data
  - Does not take very much data
  - Can use prior or expert information
  - Not all variables need to be observable
  - Can deal with missing values
- Learning model structure is hard!
  - Requires massive data
  - Or, expert input
- Discuss learning in another lesson

Focus on representation with graphical models

Observations /Evidence →

Representation
Conditional Independency

→ **Inference**
Posterior distribution
Optimal decisions

→ **Actions/ Decisions**

Learning
Structure
Parameters

Schematic of intelligent agent using directed graphical model

Some Graph Terminology: Nodes and edges
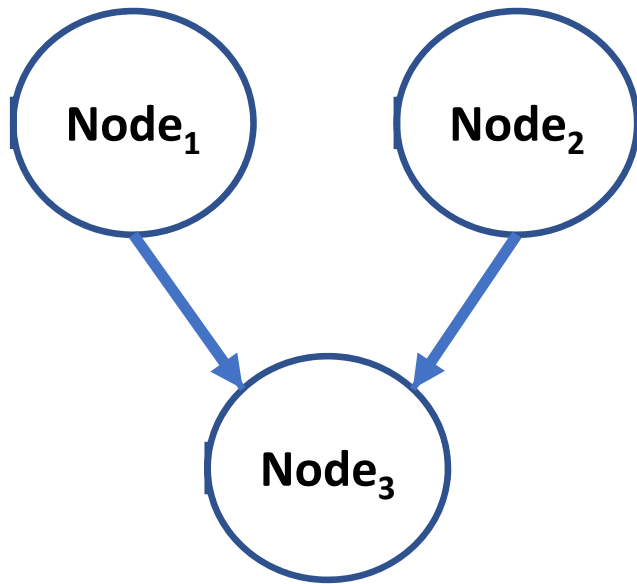
- **Nodes or vertices** contain **conditional probability distributions (CPD)**

- **Undirected edges** define the connectivity between the nodes or the **skeleton of the graph**

- Nodes connected by edges are **neighbors**

- Information or **influence** flows along **directed edges**

Some Graph Terminology: Relationships
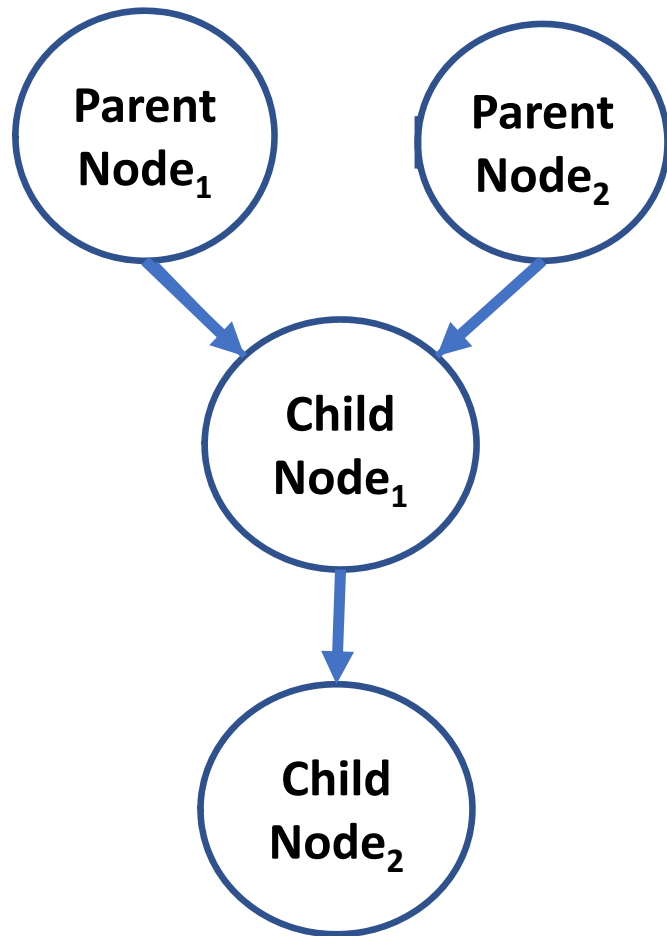
- **Parent nodes** have directed edges to, and **influence**, child nodes

- **Child nodes** are descendants of parents

- **Ancestors** are parents, grand parents, etc.

- Example; express ancestors of child node 2:

    $$PA(CN_2) = \{PN_1, PN_2, CN_1\}$$

- **Decedents** are nodes receiving influence

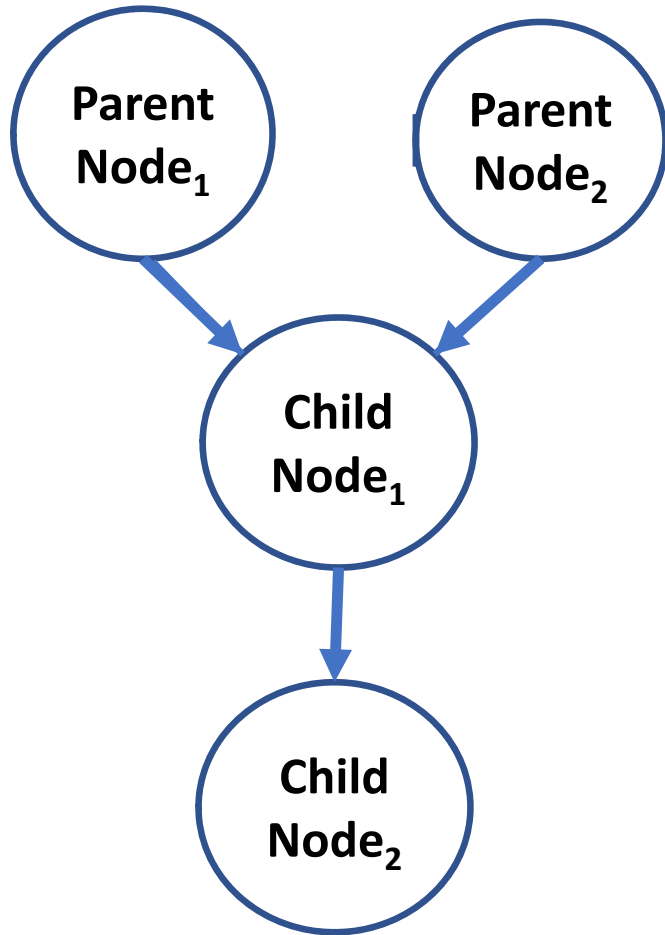- Example; express decedents of parent node 1:

    $$DE(PN_1) = \{CN_1, CN_2\}$$

Some Graph Terminology: Degree of nodes

- **Degree** of a node is the **number of neighbors**

- Example **in degree** of a child node 1 can be expressed:

  IN(CN$_1$) = 2

- Example **out degree** of a child node 1 can be expressed:

  OUT(CN$_1$) = 1

Some Graph Terminology: Special nodes

Parent Node$_1$

Parent Node$_2$

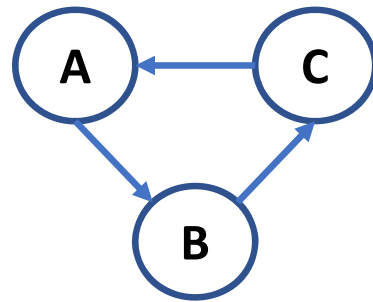Child Node$_1$

Child Node$_2$

- A **root node** has **no ancestors**
  Root nodes = {PN$_1$, PN$_2$}
- A **Leaf node** has **no children**
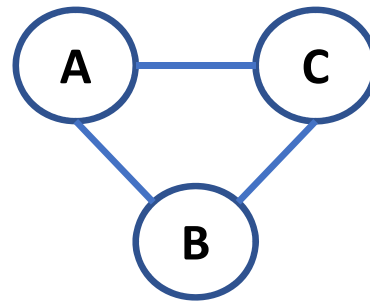 Leaf node = {CN$_2$}
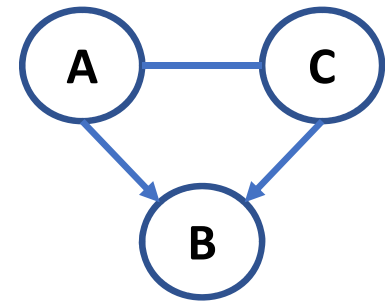
Some Graph Terminology: Types of graphs
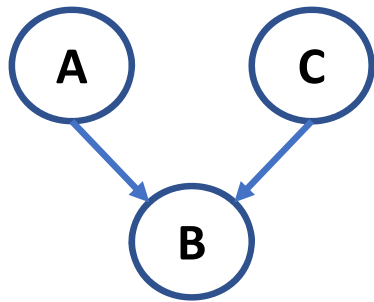
Directed Acyclic Graph

Directed Cyclic Graph

Undirected Graph

Partially Directed Graph

Independencies in Directed Graphical Models

We need to model the **independency structure of a multivariate probability distribution**

$$p(x_1, x_2, \ldots x_n)$$

- We use the notation **I(p)** to represent the independency structure of the distribution

- We use the notation **I(G)** to represent the independency structure of the graphical model

- Is it possible to have:

  I(p) = I(G)

- **In general no,** but we can still get a useful model

Independencies in Directed Graphical Models

Factorizing a distribution greatly reduces computational complexity

- A **bivariate distribution can be factored** as a conditional probability distribution and an unconditional probability distribution:

$$P(A, B) = P(A|B)P(B) = P(B|A)P(A)$$

- Notice that the **factorization is not unique**

Independencies in Directed Graphical Models

Factorizing a distribution greatly reduces computational complexity

- The **chain rule of probability** is the key to factorizing a multivariate distribution

- First, a multivariate distribution can be factored:

$$P(A_1, A_2, A_3, A_4 \dots, A_n) = P(A_1 | A_2, A_3, A_4, \dots, A_n) \, P(A_2, A_3, A_4 \dots, A_n)$$

- Continue factoring on the right hand side eventually yields:

$$= P(A_1 | A_2, A_3, A_4, \dots, A_n) \, P(A_2 | A_3, A_4 \dots, A_n) \, P(A_3 | A_4 \dots, A_n) \dots P(A_n)$$

Independencies in Directed Graphical Models

Factorizing a distribution greatly reduces computational complexity

- Factorization using the **chain rule of probability is not unique**

- For example, the multivariate distribution can be factored:

$$P(A_1, A_2, A_3, A_4 \ldots, A_n)$$
$$= P(A_n|A_{n-1}, A_{n-2}, A_{n-3}, \ldots, A_1) \, P(A_{n-1}|A_{n-2}, A_{n-3}, \ldots, A_1) \, P(A_{n-2}|A_{n-3}, \ldots, A_1)$$
$$\ldots p(A_1)$$

Independencies in Directed Graphical Models

Factorizing a distribution greatly reduces computational complexity

- Factorization can be performed in any other order, which can be expressed generally as the product of conditional distributions:

$$P(\bigcap_{k=1}^{n} A_k) = \prod_{k=1}^{n} p(A_k \mid \bigcap_{j=1}^{k-1} A_j)$$

- Unfortunately, finding the **best factorization is an NP complete problem**

Independencies in Directed Graphical Models

Factorizing a distribution greatly reduces computational complexity

- For a directed graph the **choice of parents determines the semantics**

- The factorization of the entire distribution are defined by **global semantics** of the DAG

- The global semantics are expressed:

$$P(X) = \prod_{i=1:d} P(X_i | \{parents(X_i)\})$$

Distribution Factorization Example

Example of Factorizing a distribution for probabilities that a student getting a job based on GRE score and letter of recommendation

- A student is seeking a job as a machine learning engineer
- The employer will make a hiring decision based on her GRE score and the quality of a letter of recommendation from her machine learning course professor
- The student's GRE score is only dependent on her intelligence
- The student's grade in the machine learning course is dependent on both her intelligence and the difficulty of the course
- Unfortunately, the professor is absent minded and will base her letter only on the student's grade

Distribution Factorization Example

Example of Factorizing a distribution for probabilities that a student getting a job based on GRE score and letter of recommendation

- The joint discrete distribution is: P(D, I, S, G, L)

    D = {0,1}, difficulty of the student's machine learning course

    I = {0,1}, intelligence of the student

    S = {0,1}, student's GRE score

    G = {0,1,2}, student's grade in the machine learning course

    L = {0,1}, quality of machine learning professor's recommendation letter

- The full table of the join discrete distribution is comparatively large:

    2 x 2 x 2 x 3 x 2 = 48 cases

Distribution Factorization Example

Use **conditional probability tables (CPTs)** to factor distribution
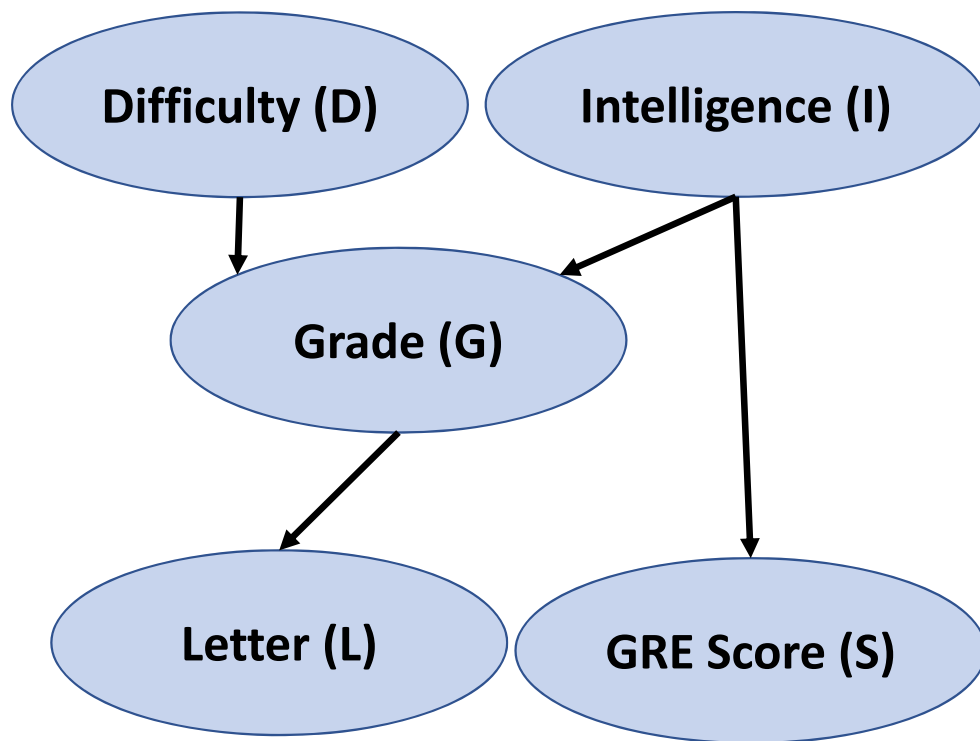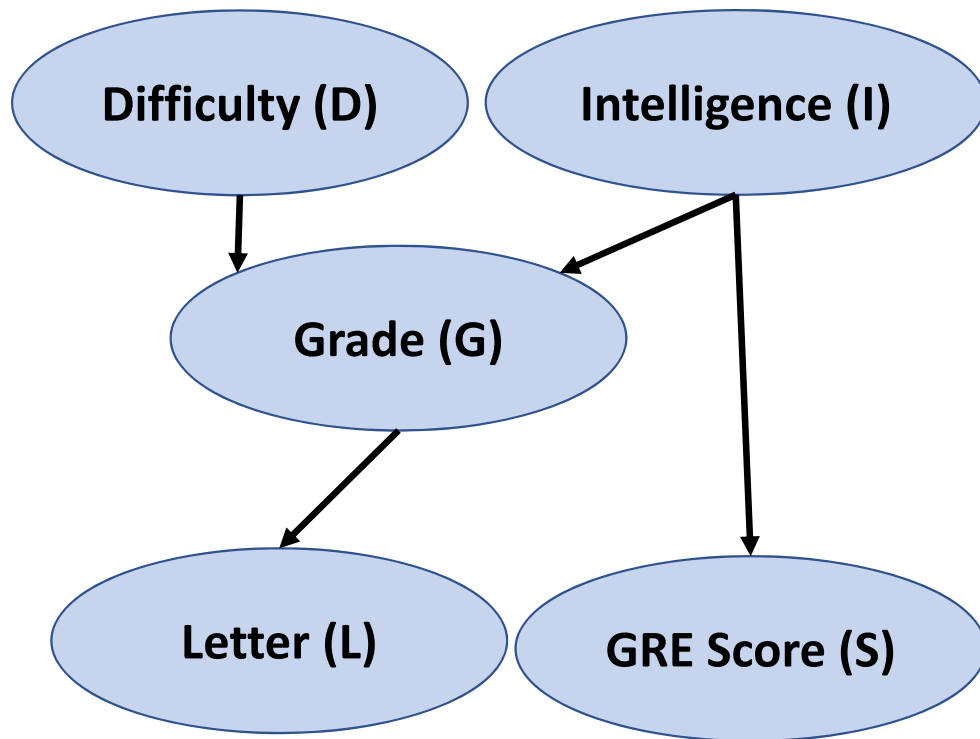

- Difficulty, D, and Intelligence, I, are independent of all other variables
- Grade, G, depends on Difficulty and Intelligence
- GRE Score, S, depends only on intelligence
- Quality of Letter, L, depends only on the grade

Distribution Factorization Example

Use **conditional probability tables (CPTs)** to factor distribution

The **semantics** of the graphical model are:

$$D \perp I, S, G. L$$
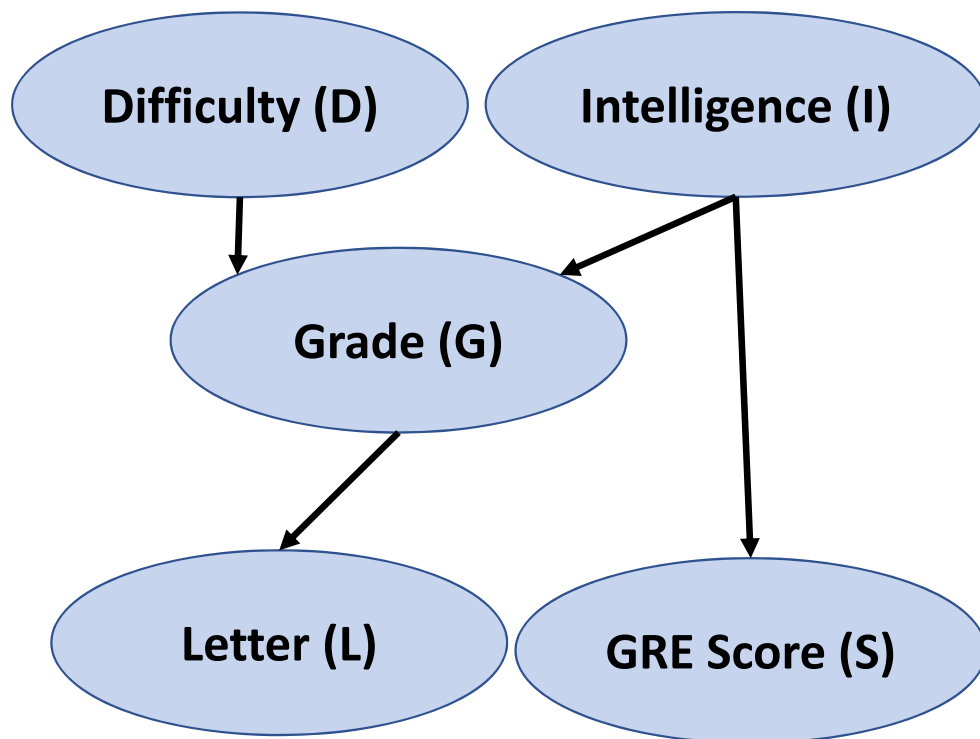
$$I \perp D, S, G. L$$

$$G \perp S, L \mid I, D$$

$$S \perp G, D, L \mid I$$

$$L \perp I, S, L \mid G$$

Distribution Factorization Example

What have we gained by factoring the distribution?

- The joint distribution can now be expressed:

$$P(D, I, S, G, L) =$$
$$P(D)\, P(I)\, P(S|I)\, P(G|I, D)\, P(L|G)$$

- One **joint table replaced with 5 CPTs**

- Largest table is now:

  2 x 2 x 3 = 12 cases

- Total cases in all tables:

  2 + 2 + 12 + 6 + 4 = 26

- For large scale problem reduction in table size can be orders of magnitude!

Independencies Structures in Directed Graphical Models

- How can we define the local directed graph structures that define independencies?

- There are only 4 types of structure that determine independency properties

- These are considered **local independencies** expressed by **local semantics**

- **Definition:** The **local semantics** of a DAG specifies that each node is conditionally independent of its non-descendants given its parents

- Understanding these local semantics enables analysis of complex graphs globally

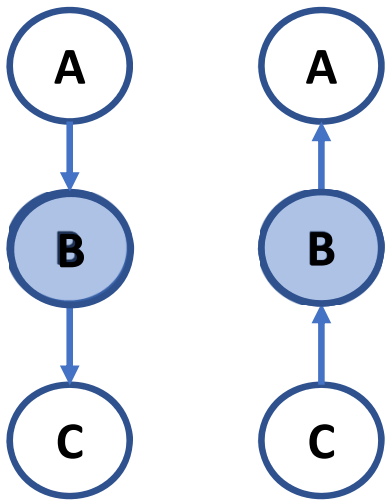# Independencies Structures in Directed Graphical Models

There are 4 types of graph structures that govern independencies
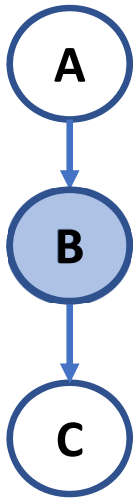
- A **cascade structure**, or **causal relationship:**
$$A \rightarrow B \rightarrow C$$
- In general A and C are not independent, $A \not\perp C$
- However, if **B is observed** then A and C are **conditionally independent,** $A \perp C \mid B$

Independencies Structures in Directed Graphical Models

There are 4 types of graph structures that govern independencies

- Reversing the arrows **does not change the independencies**
- Gives an **evidential structure:**
$$C \rightarrow B \rightarrow A$$
- In general A and C are not independent, $A \not\perp C$
- However, if **B is observed** then A and C are **conditionally independent**, $A \perp C \mid B$
- **Independencies**, I(G), for a graphical models are **not unique**

Independencies Structures in Directed Graphical Models

There are 4 types of graph structures that govern independencies

- A **common cause**, or **common parent relationship:**

$$B \to A \text{ and } B \to C$$
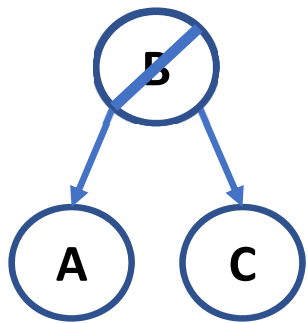
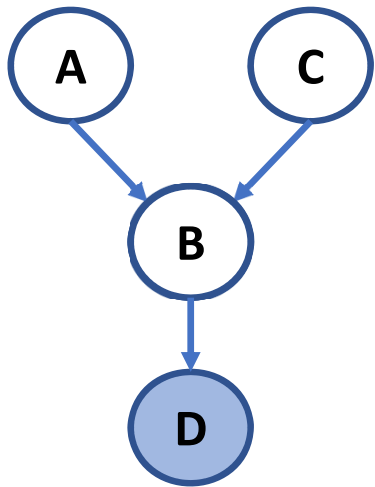- In general A and C are **independent,**

$$A \perp C \text{ or } P(A, B, C) = P(A \mid B) \, P(C \mid B) \, P(B)$$

- However, if B is observed then A and C cannot be independent, $A \not\perp C \mid B$

- If **B is marginalized out**, then:

$$\sum_{B} p(A, B, C) = p(A) \, p(B)$$

Or $A \perp B \mid marginal(B)$

Independencies Structures in Directed Graphical Models

There are 4 types of graph structures that govern independencies

- A **common evidence**, **v-structure,** or **collider**
- In general A and C are unconditionally independent**,**

$$A \perp C \ \text{ or } P(A, B, C) = P(B \mid A, C) \ P(A) \ P(C)$$

- However, if B is observed then A and C cannot be independent, $A \not\perp C \mid B$ or,

$$P(A, B, C) = P(A, C \mid B) P(B) = P(A \mid B) \ P(C \mid B) P(B)$$

- The above **indecencies applies to ancestors of the v-structure**

Independencies Structures in Directed Graphical Models

There are 4 types of graph structures that govern independencies

- A **common evidence**, **v-structure,** or **collider**

- If B is marginalized out, then:
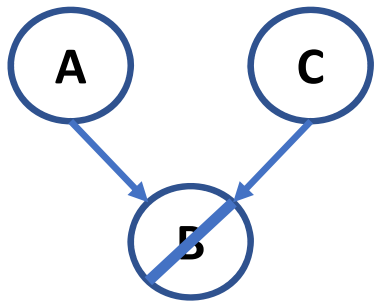
$$\sum_{B} p(A, B, C) = p(A)\, p(B)$$

Or $A \perp B \mid marginal(B)$

Mapping Independencies in Directed Graphical Models

- We have looked at simple rules for local independencies

- How do we extend these concepts to complex graphs?

- Use the concept of **D-separation**:

  > **Definition:** Given subsets *X, Y* and *Z, X* and *Y* are conditionally independent or **D-separated** conditioned on the subset *Z* if they are separated on the moralized graph.

- We can say that, *X* and *Y* are **D-separated if all paths** between them in *Z* are **blocked**

- We can state d-separation in terms of the **local Markov assumption**:

$$I(G) = \{X \perp Z \mid Y : \ dsep_G(X : Z|Y)\}$$

Note: we will discuss moralization of graphs in the next lecture

Mapping Independencies in Directed Graphical Models

Relating local semantics to global semantics
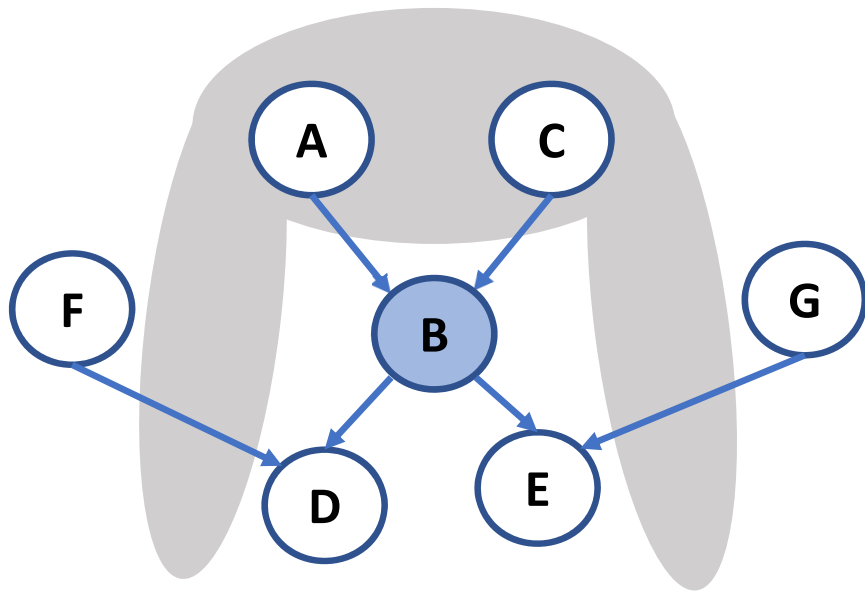
- By local semantics the a node is conditionally independent of its non-descendants.
- The local semantics map to the global semantics and vice versa

$$local\ semantics \Longleftrightarrow global\ semantics$$

Mapping Independencies in Directed Graphical Models
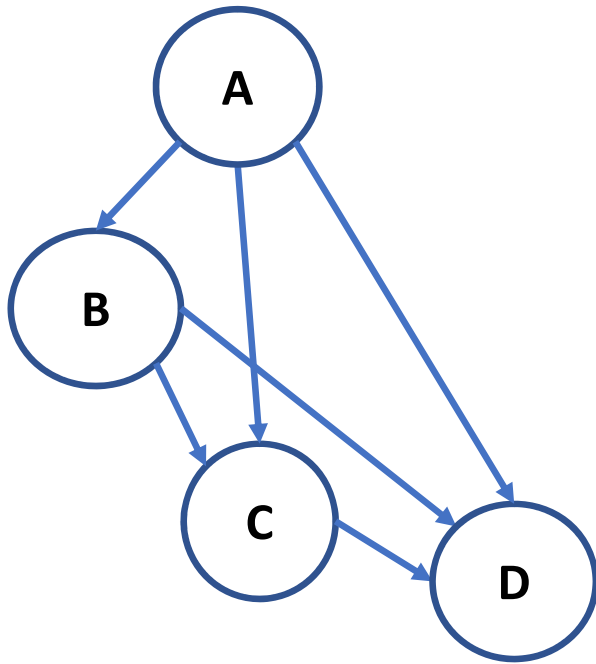
Independence map

Definition: A DAG, G, is an **independence map** or **I-map** of a distribution P if $I_l(G) \subseteq I(P)$, where $I(P)$ is the set of independencies of the distribution P and $I_l(G)$ is the set of independencies of the DAG. We can express this relationship as:

$$(X \perp Y \mid Z_G) \Rightarrow (X \perp Y \mid Z_P)$$

- This relationship is not unique and there are can be multiple graphs for which $I_l(G) \subseteq I(P)$

Mapping Independencies in Directed Graphical Models

Fully connected DAG is an I-map for any distribution

- This means we can **always create a DAG**, G, which is an **I-map of any distribution**, P, such that, $I_l(G) \subseteq I(P)$!

- But, a more compact representation is a **minimal independence map**

  Definition: A DAG, G, is a **minimal I-map** for a distribution P if removal of even a single edge renders G not an I-map.

Mapping Independencies in Directed Graphical Models

Dependency map

Definition: A graph G is a **dependency map** or **D-map** of a distribution P if the graph contains every conditional independence in P.

We can represent this relationship as:

$$(X \perp Y \mid Z_G) \Leftarrow (X \perp Y \mid Z_P)$$

Mapping Independencies in Directed Graphical Models

Perfect map

Definition: If a graph G is **both an I-map and a D-map** of a distribution P we say that G is a **prefect map** of P.

We can write this relationships as:

$$(X \perp Y \mid Z_G) \Leftrightarrow (X \perp Y \mid Z_P)$$

- It would be nice if a graph were a perfect map of a distribution
- This will rarely be the case in real world problems.
- Thus, a perfect map is mostly useful as a reference point in developing probabilistic graphical models.
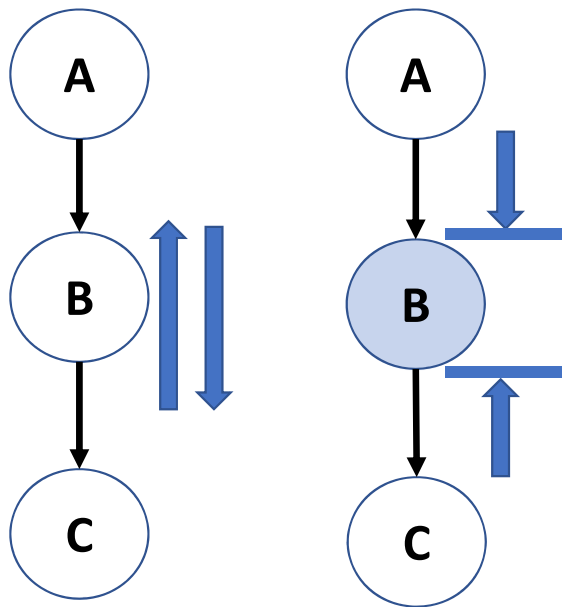
Paths and Bayes' Ball Algorithm

Two subsets of variables, X and Y are **D-separated if there is no active trail or path** though the subset Z between them

Is there a way to determine if a **path or trail is blocked or active**?

- The **Bayes' ball algorithm** determines if a path is **active**
- The ball rolls along a path
  - If the ball can complete the path, it is active
  - If the ball is blocked so is the path
- There are only a few simple rules to the Bayes' ball algorithm
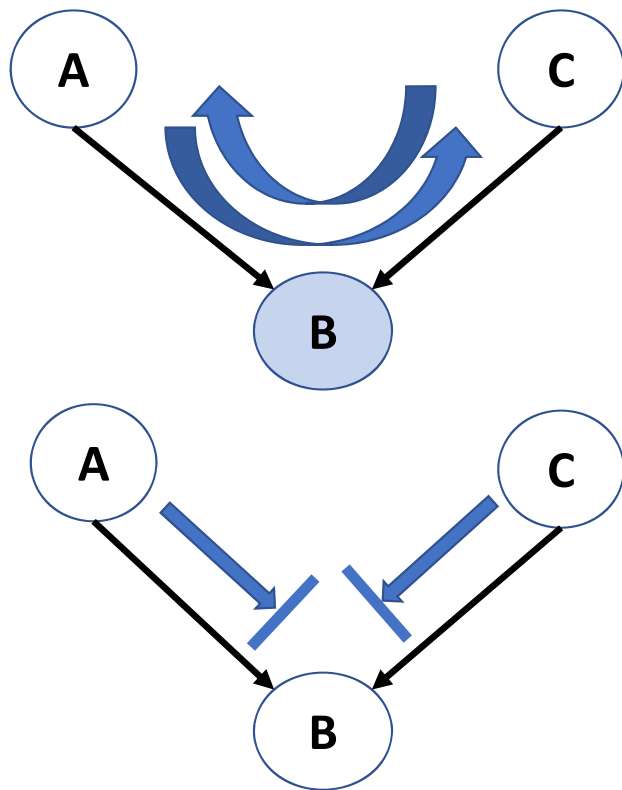
Paths and Bayes' Ball Algorithm

Rules for the Bayes' ball algorithm: **Causal trail or evidential trail**

- For a **causal relationship** or **evidential structure:**
- The ball can roll through: if **B is not observed the path is active**
- If **B is observed the path is blocked**

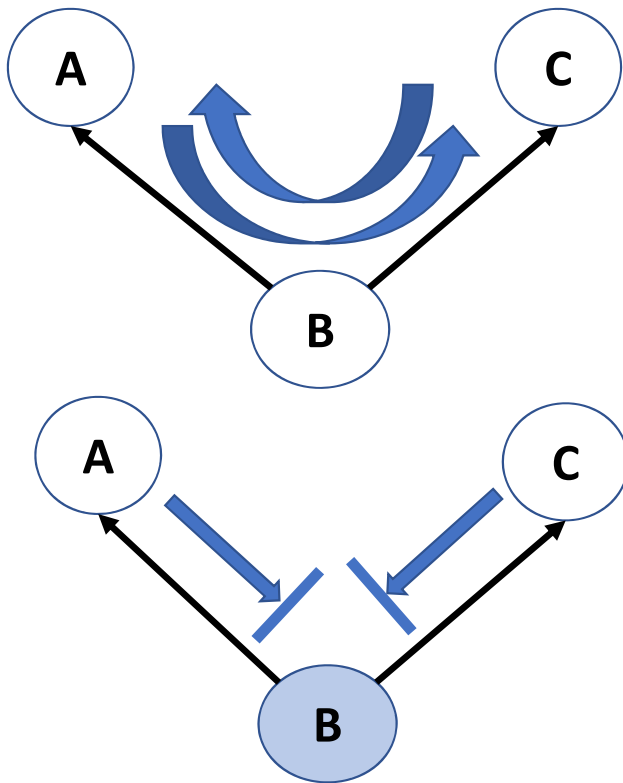# Paths and Bayes' Ball Algorithm

Rules for the Bayes' ball algorithm: **Common cause trail**

- For a **V-structure** or **common evidence:**
- The ball can roll through: if **B is observed the path is active**
- If **B not is observed the path is blocked**
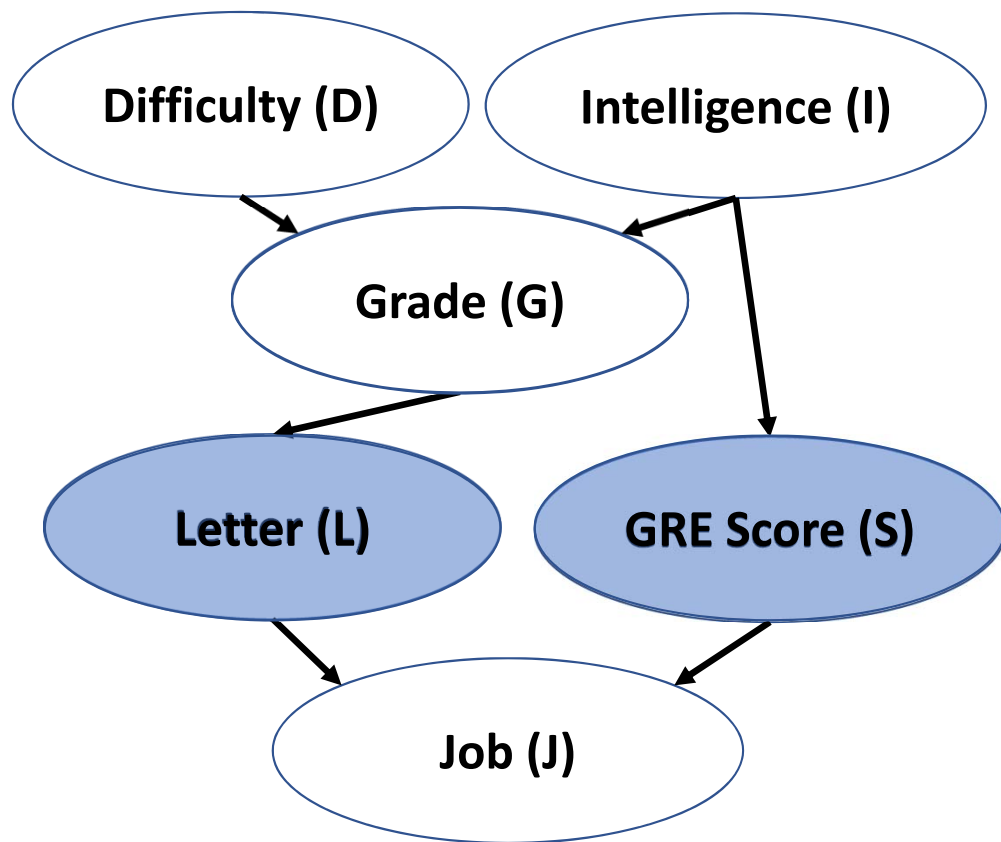
Paths and Bayes' Ball Algorithm

Rules for the Bayes' ball algorithm: **Common effect trail**

- For a **common cause** or **common parent:**
- The ball can roll through: if **B is not observed the path is blocked**
- If **B is observed the path is blocked**

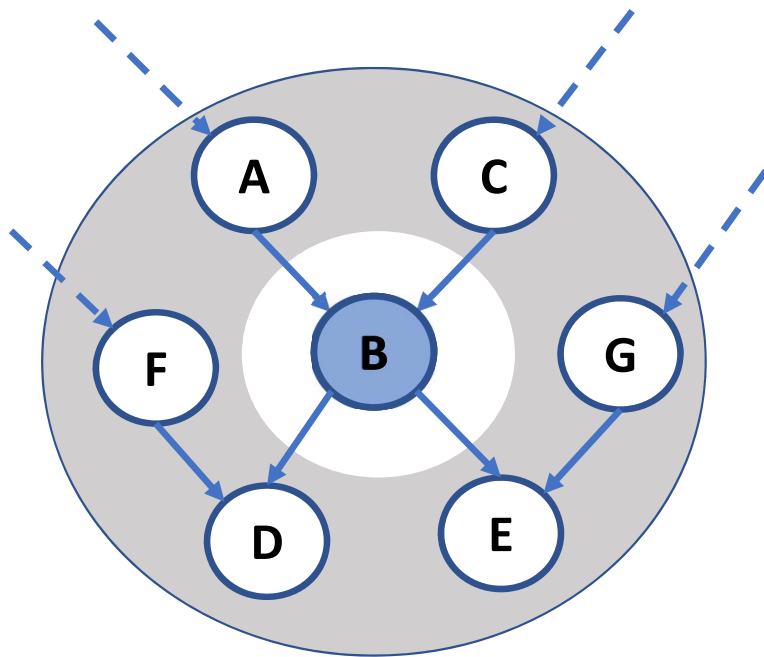# Mapping Independencies in Directed Graphical Models

Examples of D-separation on a graph

- Let X = {D}, Y = {I}, Z = {G}
- If G is not observed, X and Y are D-separated

- Or, let X = {D, I, G}, Y = {J}, and Z = {L,S}
- If L and S are observed then X and Y are D-separated

Mapping Independencies in Directed Graphical Models

For a DAG, any node is conditionally independent or all others given its **Markov Blanket**

The **Markov blanket** includes**:**

- Parents

- Children

- Parents of children

- Above define Markov blanket

Summary

In many cases, can factorize a distribution, P, or a directed graph, G

Definition: Given a distribution P and a Bayesian network G, P factorizes as a set of CDPs specified as the nodes of G

- Graph can be an I-map of P

- P can be a D-map of G

- A perfect map is both an I map and a D-map

- Useful model is rarely perfect map