

CSCI E-82a

Probabilistic Programming and AI

Introduction

Steve Elston



HARVARD
Extension School

Why Probabilistic AI?

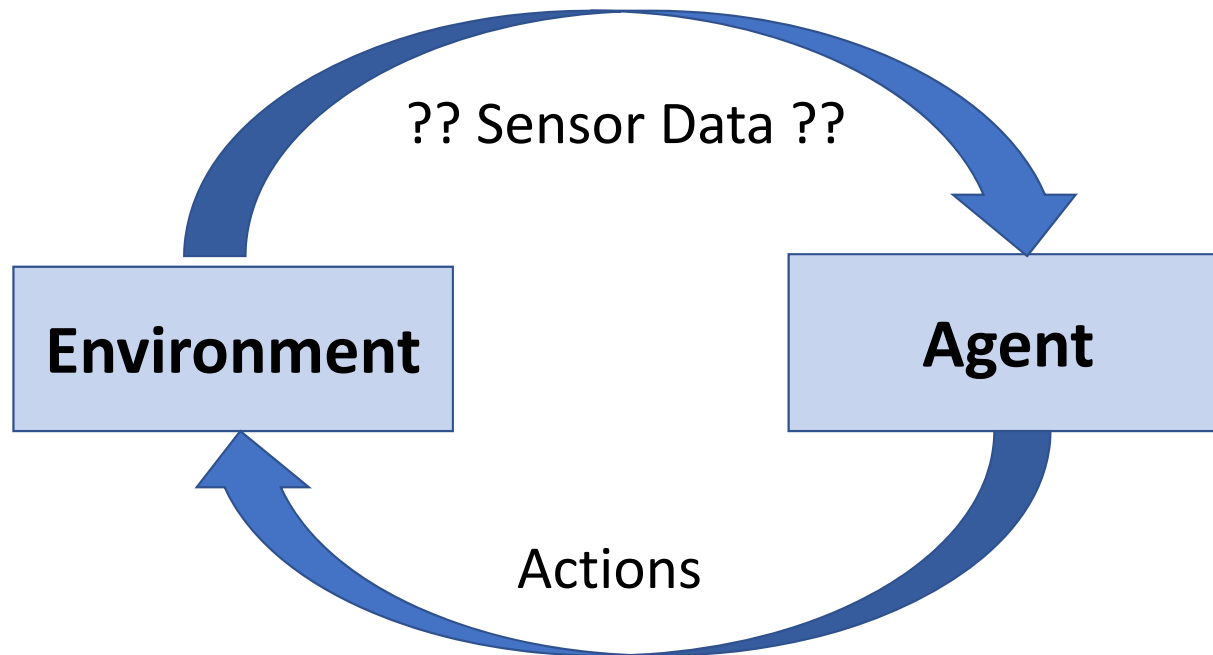
The common theme of this course is making **optimal decisions** in **complex and uncertain environments**

- Intelligent agents must interact with a complex world
- Complex environments lead to uncertainty
- Agents require algorithms that deal with uncertainty
- Probabilistic models, such as **Bayesian models** and **Markov decision processes (MDP)**, allow us to address these problems

Why Probabilistic AI?

Intelligent agent interacts with uncertain environment

- Information from the environment is incomplete and prone to errors
- Agent must take optimal actions given uncertain information



The Intelligent Agent

Fundamental functions of a probabilistic intelligent agent

- **Representation:** A good representation is often the key to good machine intelligence. A good representation is a mapping of the model and the environment. Good representation is key to effective AI!
- Representations are often **approximate** given **high complexity** of real world




Representation

The Intelligent Agent

Fundamental functions of a probabilistic intelligent agent

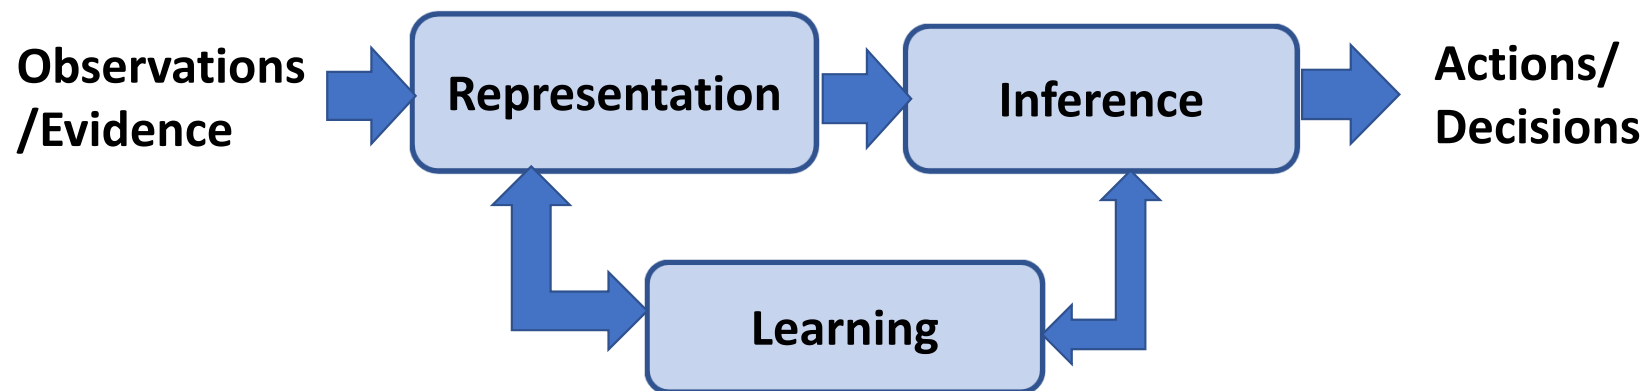
- **Inference or Reasoning:** The process of computing actions or decisions from **queries** of the model given the **evidence**. In the simplest form a query returns a mathematical result, such as the **marginal probability distribution** or the **maximum a posteriori** value.
- Reasoning computes a specific action which is applied to the environment.



The Intelligent Agent

Fundamental functions of a probabilistic intelligent agent

- **Learning:** The agent performs **learning** using data or **evidence** to update the model. The evidence is observed by **sensors** which provide information to the model on the **state of the environment**



Uncertainty in the Environment

Agent must navigate to destination

- Plans optimal route
- How much does the traffic volume change?
- Does the plan account for road repair?
- Does an accident block a route?
- In other words, which decisions are required to minimize travel time?
- Poor response to unexpected information is known as **brittleness** in a model

Uncertainty in the Environment

Integrate sensors for collision avoidance in self-driving car

- Sensors have different range and accuracy
- How are sensors affected by fog, rain or darkness?
- How accurate is traffic sign recognition?
- What is the response of each sensor to snow and ice covered roads?
- In other words, what is the posterior probability that a change in speed or direction is required?

Uncertainty in the Environment

Unobservable information adds to uncertainty

- The intentions of other drivers
- The cards held by other players in a game of poker
- The spot price of wheat in the future
- Net result is **incomplete information**

Probabilistic Reasoning Recognized as Fundamental Method

Check out videos: https://amturing.acm.org/award_winners/pearl_2658896.cfm

The screenshot displays the ACM Turing Award website. At the top, there's a header with the ACM logo and the text "A.M. TURING AWARD". Below this, a navigation bar includes "A.M. TURING CENTENARY CELEBRATION WEBCAST", a search bar, and a grid of laureate portraits. A secondary navigation bar lists "A.M. TURING AWARD LAUREATES BY...", "ALPHABETICAL LISTING", "YEAR OF THE AWARD", and "RESEARCH SUBJECT". The main content area features a large portrait of Judea Pearl on the left. To his right, his name "JUDEA PEARL" is displayed with a "DL" icon, followed by "United States – 2011". Below this, a "CITATION" section reads: "For fundamental contributions to artificial intelligence through the development of a calculus for probabilistic and causal reasoning." At the bottom, five icons represent different resources: "SHORT ANNOTATED BIBLIOGRAPHY", "ACM TURING AWARD LECTURE VIDEO", "RESEARCH SUBJECTS", "ADDITIONAL MATERIALS", and "VIDEO INTERVIEW". A "Photo-Essay" link is also visible at the bottom left.

acm
MORE ACM AWARDS

A.M. TURING AWARD

A.M. TURING CENTENARY CELEBRATION WEBCAST

A.M. TURING AWARD LAUREATES BY...

ALPHABETICAL LISTING YEAR OF THE AWARD RESEARCH SUBJECT

JUDEA PEARL DL

United States – 2011

CITATION

For fundamental contributions to artificial intelligence through the development of a calculus for probabilistic and causal reasoning.

SHORT ANNOTATED BIBLIOGRAPHY ACM TURING AWARD LECTURE VIDEO RESEARCH SUBJECTS ADDITIONAL MATERIALS VIDEO INTERVIEW

Photo-Essay

About Your Instructor

- Principle Consultant at Quantia Analytics
- Instructor, Harvard Extension School, University of Washington
- MS and PhD in Geophysics from Princeton University
- Work in machine learning starting in 1980s
- Co-founded analytics businesses
- Worked in a number of areas:
 - Capital markets risk
 - Image analysis
 - Fraud detection
 - Forecasting
 - Failure prediction

About Your Teaching Fellow

- Sarah Asano – asano.sar@gmail.com
- Electro-Optical Engineer, Lockheed Martin, Sunnyvale, California
- MS Robotics, Carnegie Mellon University
- BS Mechanical Engineering, California Institute of Technology
- Experience in:
 - App development
 - Game development
 - Internet of things
 - Embedded systems
 - Robots

About This Course

Focus on two different classes of probabilistic algorithms

- Graphical models
 - Computationally efficient method to compute posterior probabilities distributions
 - Sequential decision models
 - Explainable models
- Reinforcement learning algorithms
 - Agent learns by experience
 - Model free
 - Learn policy for complex and stochastic environment
- Models related through Markov Decision Processes (MDP)

About This Course

Grading is based on hands on work and class participation

- Homework assignments – 70%
 - Assignment most weeks
 - Focus on hands-on coding
 - **Read directions carefully and answer all questions;** don't miss points!
- On campus weekend – 30%
 - 9am – 5 pm Dec 7-8. **You must attend the entire session** for course credit!
 - Meet at one Brattle Square, Cambridge
 - Team challenges
 - Book rooms, etc. early

About This Course

Course participation

- Your participation important to get maximum value from this course!
 - Students who attend lection and sections tend to do better.
- On-line lecture – Wednesdays 5:50 – 7:50 pm US Eastern Time
 - Lecture focused on theory
 - Lectures will be recorded
 - Please remind your instructor to record!!
- Section – TBD
 - Section focused on code, questions and homework
 - Perhaps, some background supplement for theory

About This Course

Text Books

- Readings are from two text books
- Both available at the Coop: <https://tinyurl.com/300-F19-CSCI-E-82A-1>
- Or free pdf downloads
 - Bayesian Reasoning and Machine Learning, Barber, 2012, Cambridge University Press:
<http://web4.cs.ucl.ac.uk/staff/D.Barber/textbook/091117.pdf>
 - Reinforcement Learning, an introduction, Second edition, Sutton and Barto, 2018, MIT Press:
<https://mitpress.ubli.sh.com/book/reinforcement-learning-an-introduction-2>

About This Course

Other reference sources I draw material from:

- *Artificial Intelligence, A Modern Approach*, Stuart Russell and Peter Norvig, Prentice Hall, Third edition, 2010
- *Probabilistic Graphical Models, Principles and Techniques*, Daphne Koller and Nir Freedman, MIT Press, 2009
- *Decision Theory Under Uncertainty: Theory and Applications*, Kochenderfer, et. al., MIT Press, 2015.
- *Machine Learning: A Probabilistic Perspective*, Murphy, MIT Press, 2012.
- *Deep Learning*, Ian Goodfellow, Yushua Bengio, and Arron Courville, MIT Press, 2016

About This Course

Getting help with this course – essential component of class participation

1. Plan to attend the section
 - Bring your questions for class discussion
2. Use Piazza – <https://piazza.com/class#fall2019/cscie82a>
 - Ask questions
 - Answer questions
3. Email Steve – stephen.elston@quantia.com
 - Please only ask **questions of a private nature**; e.g. grading questions
 - Please direct general questions on course material and homework to the aforementioned venues – if you have a question, others likely will too!
4. Grading questions: email Sarah – asano.sar@gmail.com

Course Materials

- Obtain course materials from course Github repository
 - https://github.com/StephenElston/CSCI_E_82A_Probabilistic_Programming
 - Course material will be updated regularly – **plan on doing a pull regularly**
- Homework assignments will be at:
https://github.com/StephenElston/CSCI_E_82A_Probabilistic_Programming/Homework
- Submit completed homework and receive grades in Canvas

First Assignments

- Homework 0 – Self-paced
 - Review of probability concepts – In Github repository
 - Not graded
 - **Decide if this class is for you!**
- Homework 1 – Directed graphical models
 - Due September 18 at 24:00 (midnight) US Eastern Time